

Analyzing and Predicting Court Efficiency in India Using Machine Learning and Geospatial Visualization

Aditya Singh^a and Vikrant Mishra^a

^aCluster Innovation Centre, University of Delhi, Delhi, INDIA - 110007

ABSTRACT

The efficiency of Indian district courts is a critical aspect of judicial performance, directly impacting access to justice and public trust in the legal system. This study evaluates court efficiency using the Court Efficiency Index (CEI), derived through Principal Component Analysis (PCA), combined with key judicial metrics such as case disposal rates and pending case growth. GeoJSON-based geospatial visualization in QGIS was used to map CEI at district and state levels, revealing stark regional disparities. Additionally, a Gradient Boosting Regressor (GBR) was applied to predict future trends in judicial efficiency. The results provide actionable insights for judicial reforms and policy interventions, addressing inefficiencies and promoting equitable access to justice across India.

KEYWORDS

Court Efficiency Index, Machine Learning, QGIS, Judicial Reforms, Gradient Boosting Regressor

1. Introduction

1.1. Background

India's judiciary plays a pivotal role in sustaining democracy. However, the legal system faces significant challenges, including prolonged delays, a backlog of over 45 million cases, and unequal distribution of resources. These inefficiencies erode public trust and disproportionately affect marginalized populations. Addressing these challenges requires innovative data-driven solutions that quantify judicial performance and predict future trends to ensure equitable justice delivery.

1.2. Objectives

The objectives of this study include:

- Developing a Court Efficiency Index (CEI) using Principal Component Analysis (PCA) for quantifying judicial efficiency.
- Utilizing QGIS for geospatial mapping of efficiency across Indian districts and states.
- Employing machine learning techniques to predict future judicial efficiency trends.

These approaches aim to provide policymakers with actionable insights for reform.

1.3. Scope

This research provides a multifaceted analysis of judicial efficiency, focusing on civil, criminal, and total cases at both district and state levels. By integrating advanced data preprocessing, geospatial visualization, and predictive modeling, the study not only identifies systemic inefficiencies but also offers actionable recommendations for targeted reforms.

The scope of this research extends beyond the judiciary to broader governance and policymaking frameworks. An efficient judicial system is fundamental to upholding the rule of law, ensuring public trust, and fostering societal stability. As the adage "Justice delayed is justice denied" highlights, delays in case resolution often lead to a denial of justice, disproportionately impacting marginalized communities and eroding faith in the legal system. This study underscores the critical role of court efficiency in ensuring timely and equitable justice, which is vital for the credibility of democratic institutions.

By enabling policymakers to identify inefficiency hotspots, this research paves the way for strategic interventions such as the allocation of judicial resources, implementation of digital case management systems, and establishment of fast-track courts. Furthermore, the integration of predictive modeling with geospatial analysis equips stakeholders with the tools to anticipate future inefficiencies and take proactive measures.

The findings of this study also hold relevance for broader governance frameworks. Efficient courts not only expedite dispute resolution but also promote economic growth by creating an environment of legal certainty and reducing the cost of legal proceedings for individuals and businesses alike. For instance, timely judicial decisions can foster investor confidence, as legal delays are often cited as barriers to doing business in India.

In addition, the scope extends to improving judicial accessibility and equity. By identifying regions with low efficiency, this research emphasizes the need for targeted reforms to address disparities between urban and rural areas, ensuring that access to justice is not a privilege limited to certain regions. With robust judicial performance, vulnerable groups can seek and obtain justice more effectively, strengthening the judiciary's role as a pillar of democracy.

Ultimately, this study highlights the transformative potential of combining data-driven methodologies with judicial reforms to enhance court efficiency. By addressing systemic inefficiencies, the research aims to contribute to a judicial system that upholds the constitutional promise of "justice for all."

2. Data and Preprocessing

2.1. Dataset Description

The dataset, obtained from the NITI Ayog portal, includes:

- Pending cases categorized by duration (e.g., 01 year, 13 years).
- Monthly case disposals and delayed cases.
- Cases filed by women and senior citizens.

Additional population statistics from Census 2011 enabled per capita efficiency calculations. GeoJSON files provided geospatial data for mapping judicial metrics.

2.2. *Derived Metrics*

To capture the nuances of judicial performance, the following metrics were derived:

- **Delay Ratio:** $\text{Delay Ratio} = \frac{\text{Number of Delayed Cases}}{\text{Total Cases}}$
- **Old Pending Percentage:** $\text{Old Pending \%} = \frac{\text{Cases Pending > 5 Years}}{\text{Total Pending Cases}} \times 100$
- **Pending Case Growth:** $\text{Growth} = \frac{\text{CPC} - \text{PPC}}{\text{PPC}} \times 100$, where:
 - **CPC:** Current Pending Cases
 - **PPC:** Previous Pending Cases
- **Case Disposal Rate:** $\text{Disposal Rate} = \frac{\text{Cases Disposed}}{\text{Pending Cases}} \times 100$
- **Pending-to-Disposed Ratio:** $\text{Pending-to-Disposed Ratio} = \frac{\text{Pending Cases}}{\text{Cases Disposed}}$

2.3. *Normalization*

All metrics were normalized to a $[0, 1]$ range using Min-Max scaling:

$$\text{Normalized Value} = \frac{\text{Value} - \text{Min}(\text{Value})}{\text{Max}(\text{Value}) - \text{Min}(\text{Value})}$$

This ensured uniformity across variables and facilitated the PCA and machine learning processes.

2.4. *Dataset Preparation and Compatibility Adjustments*

The datasets for plotting contained discrepancies in naming conventions that required preprocessing for alignment:

- The GeoJSON file for districts had names in Pascal Case, while the court case file contained district court names in Upper Case. Additionally, some district court names differed from the district names in the GeoJSON file. To resolve this, the district court names were manually renamed in the Excel file to match the district names in the GeoJSON file.
- For state-level data, the GeoJSON file listed certain states with outdated names, such as "Orissa" instead of "Odisha" and "Uttaranchal" instead of "Uttarakhand." These names were updated to match the state names in the dataset, ensuring accurate mapping and visualization.

These adjustments were crucial to achieving consistent integration between the court case data and geospatial files, enabling accurate plotting of district and state-level judicial metrics.

2.5. *Geospatial Data Integration*

GeoJSON files containing district and state boundary data were integrated into the analysis. Using QGIS, geospatial layers were created to visualize judicial metrics, enabling thematic mapping of CEI values across regions.

3. Methodology

3.1. Principal Component Analysis (PCA)

Principal Component Analysis (PCA) was employed to derive the composite Court Efficiency Index (CEI) from the normalized metrics. PCA is a dimensionality reduction technique that transforms a dataset with correlated features into a set of linearly uncorrelated components called principal components. Each component captures a portion of the total variance in the data, with the first principal component accounting for the largest variance.

In this study, PCA was applied to assign weights to the derived metrics (e.g., Delay Ratio, Case Disposal Rate, Pending-to-Disposed Ratio). The steps involved in the PCA process are as follows:

- **Standardization:** Since the metrics had different ranges and units, they were normalized using Min-Max scaling to ensure uniformity before applying PCA.
- **Covariance Matrix Calculation:** A covariance matrix was computed to measure the relationships between the normalized features. This matrix captures the degree to which two metrics vary together.
- **Eigen Decomposition:** The covariance matrix was decomposed into eigenvalues and eigenvectors. The eigenvalues represent the amount of variance captured by each principal component, while the eigenvectors define the direction of the components.
- **Principal Component Selection:** The first principal component, which captures the highest variance, was selected for deriving the CEI. The eigenvector corresponding to the first principal component provided the weights for each metric.

By analyzing the eigenvector, metrics contributing positively to judicial efficiency, such as Case Disposal Rate, received higher weights. In contrast, metrics highlighting inefficiencies, such as Delay Ratio and Old Pending Percentage, were assigned negative or lower weights. These weights reflect the relative importance of each metric in influencing court efficiency.

The composite CEI for each district or state was calculated as a weighted sum of the normalized metrics:

$$\text{CEI} = w_1 \cdot X_1 + w_2 \cdot X_2 + \dots + w_n \cdot X_n$$

where w_i represents the weight derived from the eigenvector, and X_i is the normalized value of the i -th metric.

This approach allowed us to condense the multidimensional dataset into a single index that effectively summarizes court efficiency. The CEI not only highlights high-performing regions but also identifies inefficiency hotspots, providing actionable insights for judicial reforms.

3.2. Geospatial Visualization

Thematic maps were generated in QGIS to visualize the Court Efficiency Index (CEI) for civil, criminal, and total cases at district and state levels. These maps provided a clear representation of judicial performance across regions, revealing significant disparities in efficiency. Districts and states were color-coded to indicate varying levels of

efficiency, enabling intuitive interpretation of the results.

Regions with higher efficiency were represented in darker green shades, highlighting states like Kerala and Gujarat, which demonstrated effective judicial performance and backlog management. In contrast, regions with lower efficiency were depicted in lighter shades or near-white tones, with states like Bihar and Uttar Pradesh prominently showcasing systemic delays and inefficiencies in case disposal.

The visualization utilized six divisions of efficiency to categorize the CEI: very low, low, moderate, average, high, and very high. This classification facilitated a nuanced understanding of regional disparities and provided a basis for identifying areas requiring targeted judicial reforms. The maps also emphasized stark contrasts between urban and rural districts, underscoring the challenges faced by overburdened regions with dense populations. Overall, these thematic maps serve as valuable tools for analyzing and addressing judicial inefficiencies systematically.

3.3. Gradient Boosting Regressor (GBR)

Gradient Boosting Regressor (GBR) was used to predict Court Efficiency Index (CEI) values, leveraging normalized judicial metrics as input features. GBR is an ensemble learning technique that builds models sequentially, combining weak learners (typically decision trees) to minimize prediction errors. It works by iteratively optimizing a differentiable loss function and adjusting for residual errors of previous iterations.

3.3.1. Mathematics Behind GBR

The GBR algorithm minimizes the loss function by sequentially adding models to the ensemble, where each new model corrects the residuals of the previous ones. The process can be summarized as follows:

- (1) **Initialization:** Start with an initial prediction, which is usually the mean of the target variable:

$$F_0(x) = \arg \min_c \sum_{i=1}^N L(y_i, c)$$

Here, $L(y_i, c)$ is the loss function, y_i is the actual value, and c is a constant (mean of y).

- (2) **Loss Function:** In this study, Mean Squared Error (MSE) was used as the loss function to evaluate the model's performance. The MSE is defined as:

$$L(y_i, \hat{y}_i) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

where y_i are the actual values, and $\hat{y}_i$ are the predicted values.

- (3) **Pseudo-Residuals:** In each iteration m , pseudo-residuals are calculated, representing the negative gradient of the loss function with respect to the model's current predictions:

$$r_i^{(m)} = -\frac{\partial L(y_i, F_{m-1}(x_i))}{\partial F_{m-1}(x_i)}$$

These residuals quantify the errors in predictions, which the next weak learner will attempt to minimize.

- (4) **Fit Weak Learners:** A weak learner (decision tree) is trained to predict the pseudo-residuals:

$$h_m(x) \approx r_i^{(m)}$$

- (5) **Update the Model:** The predictions of the weak learner are scaled by a learning rate η and added to the ensemble:

$$F_m(x) = F_{m-1}(x) + \eta \cdot h_m(x)$$

Here, η (learning rate) controls the contribution of each weak learner and helps prevent overfitting.

- (6) **Iteration:** Repeat the process for M iterations or until the loss function converges to a minimum.

3.3.2. Hyperparameter Optimization

To improve model performance, key hyperparameters were optimized:

- **Tree Depth (d):** Determines the complexity of the weak learners. A deeper tree can capture more complex patterns but risks overfitting.
- **Learning Rate (η):** Scales the contribution of each weak learner. A smaller learning rate requires more iterations but ensures better generalization.
- **Number of Iterations (M):** Specifies the total number of weak learners. More iterations enhance model accuracy but increase computational cost.

3.3.3. Evaluation Metrics

The performance of the GBR model was evaluated using the following metrics:

- **R Score:** Indicates the proportion of variance in the target variable explained by the model:

$$R^2 = 1 - \frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{\sum_{i=1}^N (y_i - \bar{y})^2}$$

where y_i are actual values, $\hat{y}_i$ are predicted values, and $\bar{y}$ is the mean of y .

- **Mean Squared Error (MSE):** Measures the average squared difference between actual and predicted values:

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Lower MSE indicates better model performance.

3.3.4. Loss Function and Model Training

The MSE loss function ensures that the model minimizes the error in predicting CEI values. During training, the weak learners iteratively reduced the residual errors by

focusing on the most challenging predictions. The learning rate (η) played a critical role in balancing the model's performance and generalization, avoiding overfitting.

The GBR model proved highly effective, achieving an R^2 score of 0.95 on the training set and 0.93 on the validation set, demonstrating its capability to capture complex patterns in the data.

4. Results and Discussion

4.1. Feature Importance

The GBR model identified **Pending Case Growth** as the most significant predictor, followed by **Old Pending Percentage** and **Case Disposal Rate**. These findings emphasize the importance of backlog trends and case disposal efficiency in determining court performance.




Figure 1. Feature Importance derived from the Gradient Boosting Regressor model. The most significant features include Pending Case Growth, Old Pending Percentage, and Case Disposal Rate.

4.2. Model Performance

The Gradient Boosting Regressor (GBR) model displayed remarkable efficiency in predicting the Court Efficiency Index (CEI). This section discusses its performance through detailed graphs and analysis:

4.2.1. R Score Across Training and Validation Data




Figure 2. R Score Across Training and Validation Data. The high R scores of 0.95 (training) and 0.93 (validation) indicate that the model accurately captures efficiency trends while avoiding overfitting. The graph reflects consistent generalization across datasets.

Figure 2 demonstrates the R scores for the Gradient Boosting Regressor model as the training size increases. The blue line represents the training data, while the orange line corresponds to the validation data. Both curves converge at high R values, with the training R stabilizing around 0.95 and the validation R reaching 0.93.

This alignment highlights the model's robustness and generalization capability. The lack of a significant gap between the training and validation curves indicates minimal overfitting, suggesting that the model performs equally well on unseen data. The consistent increase in validation R as more data is utilized further affirms the efficacy of the Gradient Boosting Regressor in capturing underlying efficiency trends in judicial performance data.

4.2.2. Training Loss and Validation Loss Curve




Figure 3. Learning Curve: Gradient Boosting Regressor - Training Loss vs. Validation Loss. The curve illustrates the decline in Mean Squared Error (MSE) for both training and validation datasets as the training size increases.

Figure 3 presents the learning curve for the Gradient Boosting Regressor, showing the training and validation loss (Mean Squared Error, MSE) against the increasing size of the training dataset. The training loss (blue line) remains low and stable throughout, indicating the model's ability to fit the training data effectively without overfitting.

The validation loss (orange line) shows a decreasing trend as the training size increases, highlighting the model's improved performance and generalization capabilities with more training data. The convergence of training and validation loss towards the end of the curve signifies an optimized model with minimal overfitting.

The gap between the training and validation loss is small, suggesting that the model is neither underfitting nor overfitting the data. The declining validation loss with increasing data points emphasizes the importance of sufficient training data for robust model predictions. The consistent low MSE values validate the Gradient Boosting Regressor as an effective approach for predicting Court Efficiency Index values.

4.2.3. Training Loss vs. Iterations




Figure 4. Loss vs. Iterations. The graph demonstrates the training loss decreasing over successive boosting iterations, achieving a stable convergence. This indicates the models effectiveness in minimizing error during training.

Figure 4 illustrates the progression of the training loss, represented by the loss function (Mean Squared Error), across successive boosting iterations for the Gradient Boosting Regressor. The graph showcases a consistent decrease in the training loss as the number of iterations increases, highlighting the models efficiency in minimizing prediction errors during training.

Initially, the loss declines sharply, reflecting the model's rapid learning phase, where significant improvements are made to reduce errors by updating the weights in the boosting framework. As the iterations progress, the rate of loss reduction slows, indicating that the model is nearing its optimal fit to the training data.

The convergence of the training loss towards a low and stable value confirms that the model has effectively learned the underlying patterns in the dataset. This stability signifies that further iterations do not result in overfitting, as the models complexity is well-regulated by the hyperparameters, including the learning rate and maximum depth of the trees.

This graph demonstrates the Gradient Boosting Regressor's capability to iteratively refine predictions and minimize errors, thereby reinforcing its suitability for predicting the Court Efficiency Index. However, since the validation loss is not plotted here, further analysis of generalization performance would be required to ensure that the model maintains accuracy on unseen data. The consistent decrease in training loss nonetheless provides strong evidence of the model's capacity to capture the complexity of the judicial efficiency data effectively.

4.2.4. Accuracy vs. Iterations




Figure 5. Accuracy vs. Iterations. The graph demonstrates the convergence of R scores for training and validation data over iterations. The nearly identical accuracy for both datasets highlights optimal hyperparameter tuning.

The graph in Figure 5 illustrates the convergence of R^2 scores for the training and validation datasets across boosting iterations in the Gradient Boosting Regressor model. The training score (blue curve) and validation score (orange curve) both increase steadily with more iterations, eventually reaching near-identical values. This indicates that the model generalizes well without overfitting or underfitting, reflecting the optimal tuning of hyperparameters such as the learning rate, tree depth, and the number of boosting iterations.

Initially, there is a noticeable gap between the training and validation scores, as the model begins learning patterns in the training data. However, as the number of iterations increases, the gap narrows, and both curves plateau around an R^2 value of approximately 0.95 for training and 0.93 for validation. This behavior confirms that the model has effectively captured the underlying structure of the data while maintaining robustness against unseen data. The steady convergence of the two curves highlights the effectiveness of the learning rate in controlling updates during each iteration, ensuring consistent improvement in model performance.

4.2.5. Residuals Plot




Figure 6. Residuals Plot. The plot shows the residuals (differences between actual and predicted values) against the predicted values. The clustering of residuals around the zero line indicates the model’s accuracy, with minimal systematic error.

The residuals plot in Figure 6 visualizes the differences between the actual and predicted CEI values, plotted against the predicted values. Residuals close to zero (depicted by the red dashed line) signify accurate predictions, while larger deviations indicate areas where the model underperforms.

In this plot, most residuals cluster tightly around the zero line, suggesting that the Gradient Boosting Regressor has captured the majority of patterns in the dataset effectively. There are no distinct patterns or systematic biases visible, indicating that the model’s errors are randomly distributed, which aligns with the assumptions of a well-fitted model.

However, a few outliers can be observed in the higher range of predicted values (above 0.85). These points may reflect specific instances where the model struggles to predict accurately, potentially due to noise or unaccounted variables in the data. Despite these outliers, the overall distribution confirms the robustness of the model in predicting court efficiency indices across different districts and states.

4.2.6. Actual vs. Predicted Values




Figure 7. Actual vs. Predicted Values. The plot shows a strong correlation between predicted and actual CEI values. The points lie close to the perfect prediction line, confirming high predictive accuracy.

Figure 7 visualizes the relationship between the actual CEI values and the predicted values from the Gradient Boosting Regressor model. The red dashed line represents a perfect prediction line where the predicted values are equal to the actual values. The data points, represented by blue dots, are tightly clustered around this line, indicating a strong correlation and high predictive accuracy of the model.

The near alignment of the points with the red line suggests that the model effectively captures the variance in the data, with minimal deviation. This close match between actual and predicted values demonstrates the robustness of the GBR model, making it a reliable tool for predicting court efficiency. Minor deviations observed for some points could indicate areas for further fine-tuning of the model's hyperparameters or consideration of additional features to enhance performance.

Each of these evaluations demonstrates the Gradient Boosting Regressors robust performance, with high accuracy, minimal bias, and strong generalizability. These insights highlight its suitability for predicting judicial efficiency and identifying key factors influencing inefficiencies.

4.3. Geospatial Insights

The CEI maps created in QGIS provide a detailed visual representation of judicial efficiency at district and state levels. Four key maps were generated, highlighting both total cases and per capita efficiency indices. Each map categorizes states into six

divisions of efficiency, ranging from very low (lightest green) to very high (darkest green), based on their Court Efficiency Index (CEI) values.

4.3.1. Efficiency Classification

Each map divides regions into six efficiency categories:

- **Very High Efficiency (Dark Green):** Exceptional case disposal rates with minimal delays.
- **High Efficiency (Light Green):** Above-average performance with manageable case backlogs.
- **Average Efficiency (Yellow):** Balanced performance with room for improvement.
- **Moderate Efficiency (Orange):** Slight inefficiencies needing intervention.
- **Low Efficiency (Red):** Significant delays and backlogs indicating systemic challenges.
- **Very Low Efficiency (Dark Red):** Critical inefficiencies requiring urgent reforms.

These classifications, paired with color gradients, make it easy to visually interpret regional disparities and prioritize policy interventions.

4.3.2. District-Wise Total Cases Efficiency Index




Figure 8. District-Wise Total Cases Efficiency Index. The map illustrates district-level court efficiency using six efficiency divisions, with green shades representing high efficiency and lighter shades indicating inefficiencies. Urban centers like Mumbai and Chennai stand out as highly efficient.

Each district is classified into six divisions based on its Court Efficiency Index (CEI) score. The efficiency divisions highlight disparities in judicial performance. Districts in states like Kerala, Tamil Nadu, and Gujarat exhibit high efficiency, as indicated by darker green shades, reflecting superior case disposal rates and effective backlog

management. In contrast, regions in Bihar, Uttar Pradesh, and parts of Rajasthan are marked by significant inefficiencies, with lighter shades indicating very low to low efficiency. Districts in states like Maharashtra and Karnataka fall in the moderate to average efficiency range, demonstrating steady but improvable performance.

The classification into six divisions provides a granular perspective on judicial efficiency. Very low efficiency (CEI 0.250.372) signifies districts with critical inefficiencies, while low efficiency (0.3720.442) indicates below-average performance. Moderate efficiency (0.4420.489) reflects a mid-range level of judicial efficiency, and average efficiency (0.4890.516) represents districts with slightly above-average performance. High efficiency (0.5160.552) includes districts with robust performance, and very high efficiency (0.5520.803) denotes top-performing districts. Urban centers like Mumbai and Chennai stand out as highly efficient districts, while rural and under-resourced regions in states like Bihar and Uttar Pradesh face systemic challenges. These observations provide a comprehensive understanding of judicial efficiency, allowing policymakers to replicate best practices from high-performing districts and target inefficiencies in critical regions for reforms.

4.3.3. District-Wise Per Capita Total Cases Efficiency Index




Figure 9. District-Wise Per Capita Total Cases Efficiency Index. Darker green regions represent districts with high per capita efficiency, while lighter shades indicate inefficiencies relative to population.

The District-Wise Per Capita Total Cases Efficiency Index map visualizes judicial efficiency relative to the population size in each district. By normalizing judicial performance metrics per capita, the map highlights population-weighted efficiency, providing a more balanced comparison across districts with varying population densities.

Districts in regions like Kerala and Gujarat consistently exhibit high per capita efficiency, represented by darker green shades. These areas effectively utilize judicial resources and manage backlogs efficiently relative to their populations. Conversely, populous districts in Uttar Pradesh and Bihar appear in lighter green and near-white shades, underscoring the judicial inefficiencies in managing the larger caseloads of these

densely populated regions.

The classification ranges for the efficiency index are categorized into six levels: very low efficiency (00.0077), low efficiency (0.00770.0134), moderate efficiency (0.01340.0174), average efficiency (0.01740.023), high efficiency (0.0230.0327), and very high efficiency (0.0327467.9757). These gradations provide a detailed perspective on judicial performance disparities across districts.

It is important to note that some regions, such as Jammu and Kashmir, Telangana and west bengal, are missing from the map due to the unavailability of district-level population data in the Census of 2011. Additionally, certain districts are not represented as they were formed after 2011 and thus lack census data for population normalization. This data gap highlights a limitation in accurately assessing judicial efficiency for these regions.

The contrast between urban and rural districts is evident in this map. Urban centers such as Bengaluru and Pune, which demonstrate a combination of population density and efficient case management, exhibit high per capita efficiency. These districts serve as benchmarks for underperforming regions. On the other hand, rural districts in overburdened states require targeted reforms to enhance their judicial efficiency. This visualization provides actionable insights for addressing population-weighted disparities and optimizing judicial resource allocation across India.

4.3.4. State-Wise Total Cases Efficiency Index




Figure 10. State-Wise Total Cases Efficiency Index. Green shades represent high-efficiency states, while lighter shades indicate inefficiencies in managing court cases at the state level.

The State-Wise Total Cases Efficiency Index map offers a comparative perspective on judicial performance across Indian states. These divisions provide a macroscopic view of how effectively states manage total court cases, offering insights into systemic efficiency at the state level.

High-performing states like Kerala, Gujarat, and the northeastern states such as Sikkim and Mizoram are highlighted in darker green shades, reflecting superior ju-

dicial efficiency characterized by effective case disposal rates and minimal delays. In contrast, states such as Uttar Pradesh and Bihar are seen in lighter shades, representing systemic inefficiencies in case management and significant backlog challenges. These inefficiencies are often attributed to high case volumes, inadequate judicial infrastructure, and delays in case processing.

The classifications range from very low efficiency (0.68460.7541) to very high efficiency (0.80160.8457). This clear gradation underscores the stark contrasts in judicial performance across states, revealing regional disparities. Urbanized states with better infrastructure and digitization efforts, such as Gujarat and Karnataka, demonstrate higher efficiency. On the other hand, densely populated states like Uttar Pradesh struggle with a substantial case backlog, resulting in lower efficiency scores. This visualization serves as a critical tool for identifying priority areas for judicial reforms and resource allocation at the state level.

4.3.5. Per Capita State Judicial Efficiency Index




Figure 11. Per Capita State Judicial Efficiency Index. The map highlights the efficiency of judicial performance on a per capita basis, with green shades representing higher efficiency levels. Note: Data for Jammu and Kashmir was unavailable, and hence it is excluded from the map.

The Per Capita State Judicial Efficiency Index map provides a nuanced view of judicial performance by normalizing efficiency metrics on a per capita basis. This representation allows for a fair comparison between states with varying population densities.

States like Kerala and Himachal Pradesh, depicted in darker green shades, exhibit high efficiency on a per capita basis. This performance reflects their effective case management and judicial resource allocation relative to their smaller population sizes. In contrast, states like Bihar and Uttar Pradesh appear in lighter green shades, underscoring their struggles with judicial inefficiencies despite having large populations and higher case volumes.

A key note for this visualization is the absence of data for Jammu and Kashmir due to unavailability of Census data. As a result, the region is not represented on the

map. The classification ranges from very low efficiency (0.00070.0068) to very high efficiency (0.035467.9757), offering a clear gradation that highlights the disparities in judicial performance. The map serves as a critical tool for identifying states where judicial resources need to be scaled to enhance access to justice.

4.3.6. Regional Court Efficiency




Figure 12. Regional Court Efficiency Heatmap. The visualization highlights percentage differences in judicial efficiency compared to regional averages. Green indicates above-average performance, while red signifies below-average efficiency.

India was divided into five regionsEast, North, Northeast, South, and Westbased on geographic proximity and socio-economic similarities to analyze judicial efficiency. This classification enabled meaningful comparisons of court performance across states within a regional context. The heatmap in Figure 12 illustrates the percentage differences in judicial efficiency compared to regional averages, with green shades indicating above-average performance and red shades highlighting below-average efficiency.

The analysis reveals significant disparities across regions and states. In the North, Sikkim emerged as the top-performing state, exceeding the regional average by 8.9%, reflecting its strong case disposal rates and efficient judicial processes. In stark contrast, Chandigarh showed the lowest performance, falling 11.8% below the regional average due to systemic inefficiencies. The Northeast region maintained steady efficiency, with Tripura (+4.7%) and Mizoram (+0.7%) performing above average, showcasing effective judicial management.

In the South, states like Tamil Nadu (+3.2%) and Telangana (+1.6%) demonstrated

consistently strong efficiency, aligning with the overall trend of robust judicial systems in this region. Karnataka and Maharashtra performed near their regional averages, indicating balanced efficiency without significant outliers. Conversely, in the East, states like Bihar (-6.2%) and Odisha (-4.0%) struggled with judicial inefficiencies, driven by delays and backlogs. West Bengal also underperformed, with efficiency levels 4.4% below the regional average.

The Western region showcased notable disparities, with Gujarat performing exceptionally well at +6.8% above the regional average, whereas Rajasthan fell slightly below at -2.3%. Uttar Pradesh in the East displayed a mixed performance, with some urban courts achieving efficiency improvements, while rural areas continued to struggle.

This regional analysis highlights the stark differences in judicial efficiency across India, providing a clear picture of which states are leading in performance and which require immediate attention. States like Gujarat, Sikkim, and Tripura can serve as benchmarks for judicial reforms, while underperforming states like Chandigarh, Bihar, and Odisha call for targeted interventions to address systemic inefficiencies. The heatmap effectively aids policymakers and stakeholders in understanding regional trends and prioritizing judicial resources for maximum impact.

4.3.7. State Efficiency Compared to National Average




Figure 13. State Efficiency Compared to National Average. Green bars represent states performing above the national average, while red bars indicate below-average performance. The chart highlights significant disparities in judicial efficiency across India.

The bar chart in Figure 13 compares the judicial efficiency of Indian states relative to the national average. States with green bars exhibit higher efficiency, handling cases faster than the national average, while states with red bars demonstrate lower efficiency, indicating delays and backlogs in case management.

The analysis highlights significant variations across states. Sikkim emerges as the top performer, exceeding the national average by approximately 8.5%. This performance

reflects robust judicial systems and consistent case disposal rates. Gujarat and Tripura follow closely, showcasing strong case management and efficient utilization of judicial resources. States like Kerala, Tamil Nadu, and Telangana also appear in the high-efficiency category, maintaining their reputation for effective judicial processes.

On the other end, Chandigarh ranks the lowest, with efficiency levels approximately 11.8% below the national average. This underperformance can be attributed to systemic inefficiencies, resource constraints, and delayed case disposals. Uttar Pradesh and Bihar also exhibit significant inefficiencies, with backlogs and a slow case resolution process contributing to their negative performance. West Bengal and Odisha, though not as low as Chandigarh, still fall considerably below the national average.

The chart provides a visual representation of how states align with or deviate from the national judicial efficiency benchmark. The length of the bars indicates the degree of deviation, with longer green bars representing states excelling in judicial efficiency and longer red bars highlighting areas requiring immediate intervention. This analysis underscores the need for tailored reforms in underperforming states, while top-performing states can serve as models for best practices in judicial administration.

4.4. Predictive Insights

The Gradient Boosting Regressor (GBR) model, trained on judicial efficiency metrics, demonstrates significant potential to identify regions and districts that require focused interventions. By analyzing the CEI predictions, policymakers can prioritize underperforming areas, ensuring targeted resource allocation and infrastructural improvements.

For instance, the model can highlight rural districts in states like Uttar Pradesh and Bihar where judicial inefficiency is prevalent due to high pending case growth and low disposal rates. Urban centers like Mumbai and Chennai, which consistently exhibit high efficiency, can be used as benchmarks for identifying effective practices. The model's predictive insights can guide state and district authorities to implement corrective measures, such as introducing fast-track courts, deploying additional judicial officers, or adopting digital solutions to improve efficiency.

The model's evaluation metrics validate its reliability in identifying these trends. With a Train RMSE of 0.0100, Train R of 0.9535, and a Test RMSE of 0.0069, the model achieves a Test R of 0.9681, indicating high accuracy and minimal error in predictions. Furthermore, cross-validation results reinforce its robustness, with R scores of [0.9432981, 0.93821809, 0.85145642, 0.96535269, 0.92083349] and a mean cross-validation R of 0.9238. These metrics confirm the model's generalization capability across different datasets.

By leveraging these predictions, judicial authorities can anticipate future inefficiencies and allocate resources more effectively. For example, districts with low predicted efficiency could be prioritized for introducing alternative dispute resolution mechanisms, enhancing case tracking systems, or increasing the judge-to-population ratio. State governments could use these insights to identify systemic issues, enabling reforms tailored to local challenges.

While this study primarily focuses on Visualization, its predictive potential provides a roadmap for future applications. Integrating such predictive analytics into judicial planning would ensure a data-driven approach to improving efficiency, reducing case backlogs, and fostering equitable access to justice across India.

5. Conclusion

The analysis and prediction of court efficiency in India using geospatial visualization and machine learning techniques provide actionable insights into the judicial system’s performance. The study utilized the Court Efficiency Index (CEI) derived from Principal Component Analysis (PCA) to assess judicial efficiency across districts and states. By integrating detailed metrics, geospatial mapping, and predictive modeling, the study delivers a robust framework for evaluating and improving court efficiency.

5.1. Insights from Geospatial Visualization

The geospatial maps created for districts and states revealed stark regional disparities in court efficiency. States like Gujarat, Kerala, and Sikkim demonstrated high efficiency, reflected in darker green shades on the maps, while states such as Bihar and Uttar Pradesh consistently underperformed, as indicated by lighter green or white regions. The district-wise per capita efficiency index highlighted the disparity between urban and rural areas, with urban centers like Mumbai, Chennai, and Bengaluru excelling in case management and backlog reduction.

These maps also emphasize the role of population density and judicial resource allocation, particularly in overburdened districts where efficiency is low. Missing data for regions like Jammu and Kashmir and Telangana (due to census or boundary inconsistencies) underscores the need for standardized and updated datasets for nationwide assessments.

5.2. Model Performance and Validation

The Gradient Boosting Regressor (GBR) model demonstrated exceptional performance in predicting court efficiency. With a training R of 95.35%, a test R of 96.81%, and a mean cross-validation R of 92.38%, the model showed high accuracy and robustness in capturing the relationships between judicial metrics and efficiency. Validation curves confirmed minimal bias and variance, while residual plots and accuracy vs. iterations graphs highlighted the model’s strong generalization capability.

The loss vs. iterations graph further illustrated how the model effectively minimized error, achieving stable convergence. Actual vs. predicted value plots reaffirmed the model’s high precision, with predictions closely aligning with actual CEI values. These results validate the potential of machine learning models to analyze complex datasets and provide reliable predictions for judicial efficiency.

5.3. Implications of Predictive Insights

Although the model has not yet been employed for real-world predictions, its demonstrated accuracy and reliability suggest immense potential. The trained model can identify regions requiring urgent intervention by predicting low-efficiency areas, enabling policymakers to focus on districts and states that struggle with delays and backlogs. Such insights could guide resource allocation, digitalization efforts, and targeted reforms in the judicial system.

For instance:

- The model can highlight rural districts in Uttar Pradesh and Bihar where efficiency is consistently low, suggesting the need for infrastructural upgrades and

- additional judicial resources.
- Urban centers like Bengaluru and Chennai, which consistently perform well, can serve as benchmarks for struggling regions.

5.4. *Applications and Future Scope*

This study provides a data-driven approach to evaluating and predicting judicial performance, which can be extended in multiple ways:

- (1) **Policy Reforms:** Insights from the analysis can guide reforms aimed at equitable resource distribution and backlog reduction.
- (2) **Resource Allocation:** By identifying high-need areas, judicial resources can be allocated more effectively to improve performance in underperforming regions.
- (3) **Digitalization and Automation:** Courts can leverage technology to enhance case management, drawing lessons from high-efficiency regions.
- (4) **Future Predictions:** With further refinement, the model could predict court efficiency trends over time, enabling proactive interventions.

5.5. *Limitations*

While the study provides valuable insights, certain limitations need to be addressed:

- Missing data for regions like Jammu and Kashmir and discrepancies in district and state naming conventions limited the scope of analysis.
- The models predictions are based on historical data, which may not capture sudden changes or reforms implemented in the judicial system.

This study highlights the power of combining geospatial tools and machine learning to address systemic inefficiencies in the judicial system. By identifying high-performing regions and inefficiency hotspots, the research lays a foundation for data-driven judicial reforms. The results underscore the need for continued efforts in improving court efficiency to ensure equitable access to justice across India.

6. **Acknowledgment**

We thank NITI Ayog for providing us with the dataset of the court and the QGIS team for their geospatial tools, which were instrumental in this research.

We express our heartfelt gratitude to Dr. Shobha Bagai and Dr. Harender Pal Singh for their invaluable guidance, mentorship, and support throughout the development of this research article. Their insights and encouragement were pivotal in shaping the direction and quality of this work.

We also extend our thanks to our friends and colleagues for their constant support, constructive feedback, and collaborative efforts, which greatly enriched the research process.

7. **Appendices**

For detailed implementation code, model architecture, and loss functions related to this project, please refer to our GitHub repository:

<https://github.com/Vikrantmishra/Court-Cases>
The dataset used from the NITI Ayog website:
<https://ndap.niti.gov.in/dataset/7150>

References

- [1] Bani-Taha, O. I., & Shafiq, M. O. (2020). Combining the richness of GIS techniques with visualization tools to better understand the spatial distribution of data - a case study of Chicago City crime analysis. *International Journal of Big Data Intelligence*, 7(1), 29–46. <https://doi.org/10.1504/IJBDI.2020.106177>.
- [2] Joshi, N., Shaikh, S., Tripathy, A., Sen, S., Janson, J., Karnik, S., & Varghese, B. (2019). Crime Anatomization Using QGIS. In *Proceedings of the 2019 IEEE 5th International Conference for Convergence in Technology (I2CT)* (pp. 1–4). <https://doi.org/10.1109/I2CT45611.2019.9033910>.
- [3] Kowsrihawatt, K., Vateekul, P., & Boonkwan, P. (2018). Predicting Judicial Decisions of Criminal Cases from Thai Supreme Court Using Bi-directional GRU with Attention Mechanism. In *Proceedings of the 2018 5th Asian Conference on Defense Technology (ACDT)* (pp. 50–55). <https://doi.org/10.1109/ACDT.2018.8592948>.
- [4] Li, X., Li, W., & Xu, Y. (2018). Human Age Prediction Based on DNA Methylation Using a Gradient Boosting Regressor. *Genes*, 9(9), Article 424. <https://doi.org/10.3390/genes9090424>.
- [5] dos Santos, P. T. C., Henrique, F., Garcia, V., Ferreira, V. R. S., dos Santos Neto, A. C., Souza, J. C., Manfredini, C., Frana, J. V. F., Boaro, J. M. C., Braz Junior, G., Silva, A. C., Paiva, A. C., de Oliveira, M. S. L., Alves, E. W. B. A. L., Moreira e Moraes, R. U., & Sobral Neto, J. S. (2020). Multiclass Legal Judgment Outcome Prediction for Consumer Lawsuits using XGBoost and TPE. In *Proceedings of the 2020 IEEE International Conference on Systems, Man, and Cybernetics (SMC)* (pp. 881–886). <https://doi.org/10.1109/SMC42975.2020.9283094>.
- [6] Chemin, M. (2020). Judicial Efficiency and Firm Productivity: Evidence from a World Database of Judicial Reforms. *The Review of Economics and Statistics*, 102(1), 49–64. https://doi.org/10.1162/rest_a_00799.
- [7] Fauvrelle, T. A., & Almeida, A. T. C. (2018). Determinants of Judicial Efficiency Change: Evidence from Brazil. *Review of Law & Economics*, 14(1), 20170004. <https://doi.org/10.1515/rle-2017-0004>.
- [8] Cai, J., Xu, K., Zhu, Y., Hu, F., & Li, L. (2020). Prediction and analysis of net ecosystem carbon exchange based on gradient boosting regression and random forest. *Applied Energy*, 262, 114566. <https://doi.org/10.1016/j.apenergy.2020.114566>.